

How to Simulate a Self-Driving Car

Inspired by YouTuber [Siraj Raval](#).

Initialize

- This workshop is for aspiring and active developers
- Please let's help each other out
- Ask who is a Python beginner/comfortable/pro? (show hands)
- Ask minority to disperse
- Feel free to ask questions at any time.
- Including, perhaps even especially, among yourselves.

Start with a demo



Three step process

- 1) Generate Data

- 2) Train model

- 3) Test

1) Data Generation

NVIDIA trained an [autonomous car](#) in real life.

Udacity built a [self-driving car simulator](#) to perform the same task.

1) Data Generation


NVIDIA trained an autonomous car in real life.

Udacity built a self-driving car simulator to perform the same task.


THEY DID THE EXACT SAME THING WE ARE GOING TO DO.

- Three cameras to the head of the car
- Human driver
- Collect data (i.e. steering angle, speed, throttle, brake and images)
- Train a neural net
- Drive car using a single camera

The three cameras



The three cameras



The three cameras



left



center



right

...steers you back on course after minor model errors

2) Training

machine learning approach = behavioural cloning.

- NVIDIA = 72 hours all kinds of conditions
- We will do a few laps on a single circuit.


2) Training

- steering angle
- throttle
- brake
- speed



- 'how to drive the car'

2) Training



This requires a lot of training data so we are going to simplify..

.. and just train the **steering angle**.

2) Training

label images kinda like this..



Left a bit



Straight on



Right

2) Training


This is an example of a supervised learning algorithm

Reduce the difference between:

- 1) what human did
- 2) what the computer thinks it should do


2) Training

This is an example of a **supervised learning** algorithm



2) Training


This is an example of a **supervised learning** algorithm



GIVEN A NEW IMAGE I HAVEN'T SEEN YET, HOW SHOULD I DRIVE?

3) Testing

Drive the car using the neural net.




3) Testing

We will communicate using a **server-client** architecture



3) Testing

server-client architecture



3) Testing

server-client architecture





Fork, clone and cd into [source code repo](#)



Download Self-Driving Car Simulator

Welcome to Udacity's Self-Driving Car Simulator

This simulator was built for [Udacity's Self-Driving Car Nanodegree](#), to teach students how to build self-driving courses using deep learning. See more [project details here](#).

All the assets in this repository require Unity. Please follow the instructions below for the

Available Game Builds (Precompiled builds of the simulator)


Instructions: Download the zip file, extract it and run the executable file.

Version 2, 2/07/17

[Linux Mac Windows](#)

Version 1, 12/09/16

[Linux Mac Windows 32 Windows 64](#)



Install [miniconda](#)



Install dependencies

```
conda env create -f environments.yml  
source activate car-behavioral-cloning
```

Install TensorFlow

```
conda install tensorflow
```

ALVIN (1992)



[Courtesy of Dean Fournier]

NVIDIA (2016)



Generate Data

- Run car simulator in TRAINING MODE
- Press 'r' key to select fixed output directory
- Press 'r' again to record driving skillz
- Complete between 3 to 5 laps.

Generate Data

Inspect

- Images (IMG/)
- [driving_log.csv](#)

Train the model

```
python model.py --help
```

The Model

$$y = \text{brains}(X)$$

The Model

$$y = \text{brains}(X)$$

↑ ↑
LABEL INPUT

(ALL MACHINE LEARNING MODELS EVER)

The Model

$y = \text{brains}(X)$

(ALL MACHINE LEARNING MODELS EVER)

↑
LABEL

↑
INPUT

HOW TO
DRIVE

IMAGES

(OUR MODEL)

The Model

$$y = \text{brains}(X)$$

↑
LABEL


↑
INPUT

(ALL MACHINE LEARNING MODELS EVER)

HOW TO
DRIVE

IMAGES

(OUR MODEL)



BUT WHAT ARE BRAINS?

The Model

$$y = \text{brains}(X)$$

↑
LABEL


↑
INPUT

(ALL MACHINE LEARNING MODELS EVER)

HOW TO
DRIVE


IMAGES

(OUR MODEL)



BUT WHAT ARE BRAINS?

a way of converting some input into something else
that has meaning for humans



= cat

The Model

$$y = \text{brains}(X)$$

↑
LABEL


↑
INPUT

(ALL MACHINE LEARNING MODELS EVER)

HOW TO
DRIVE

IMAGES

(OUR MODEL)




BUT WHAT ARE BRAINS?

a way of converting some input into something else
that has meaning for humans

also = arrangement of mathematical functions

Convolutional Neural Network (CNN)

- Keras
 - python machine learning library
- everything arranged in layers
 - input is first layer
 - output is last layer
- Model specifics
 - 2D Convolutional layers
 - Dropout layer
 - Fully connected layers



What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

<u>1</u>	<u>0</u>	<u>2</u>	<u>1</u>	<u>2</u>	<u>0</u>
<u>0</u>	<u>2</u>	<u>1</u>	<u>2</u>	<u>2</u>	<u>1</u>
<u>1</u>	<u>1</u>	<u>0</u>	<u>2</u>	<u>3</u>	<u>4</u>
<u>4</u>	<u>5</u>	<u>2</u>	<u>1</u>	<u>2</u>	<u>2</u>
<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>2</u>
<u>3</u>	<u>2</u>	<u>1</u>	<u>0</u>	<u>4</u>	<u>2</u>

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

0	1	0
1	1	1
0	1	0

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

0	1	0
1	1	1
0	1	0

= 4

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2

What is a 2D convolution layer?

- image input
- filter (a.k.a. kernel)
- dot product
- move along
- repeat
- outputs a new layer

1	0	2	1	2	0
0	2	1	2	2	1
1	1	0	2	3	4
4	5	2	1	2	2
1	0	0	0	1	2
3	2	1	0	4	2


4	7	8	10
9	6	8	13
12	8	7	9
8	3	2	9

What is a 2D convolution layer?

- Neurons are only locally connected (i.e. in width and height)
- Less computationally expensive (images too large for full connections)
- Can find spatial patterns
- Much better at generalized learning (fully connected overfit)

What is a **fully connected** layer?

- More traditional neural network
- Every single neuron in a layer is connected to every other neuron



Code time!



Code time!



Sequential is a linear stack of layers

```
from keras.models import Sequential
```

this creates a model instance

```
model = Sequential()
```

For an 8-bit image (i.e. from 0 to 255) this will normalize values between -1 to 1.

```
model.add(Lambda(lambda x: x/127.5-1.0, input_shape=INPUT_SHAPE))
```

Code time!



Sequential is a linear stack of layers

```
from keras.models import Sequential
```

this creates a model instance

```
model = Sequential()
```

For an 8-bit image (i.e. from 0 to 255) this will normalize values between -1 to 1.

```
model.add(Lambda(lambda x: x/127.5-1.0, input_shape=INPUT_SHAPE))
```



must define input shape
in the first layer

Code time!



```
model.add(Conv2D(24, 5, 5, activation='elu', subsample=(2, 2)))
```



Translation to English..

A 2D convolutional layer with 24 sets of random weights and a 5 by 5 pixel kernel.
Neurons will fire based on an exponential linear unit activation function.
The kernel will skip 2 pixels vertically and horizontally during the convolution.

Code time!



```
model.add(Conv2D(24, 5, 5, activation='elu', subsample=(2, 2)))  
model.add(Conv2D(36, 5, 5, activation='elu', subsample=(2, 2)))  
model.add(Conv2D(48, 5, 5, activation='elu', subsample=(2, 2)))  
model.add(Conv2D(64, 3, 3, activation='elu'))  
model.add(Conv2D(64, 3, 3, activation='elu'))
```



A bunch of convolutional layers

Code time!



```
model.add(Dropout(args.keep_prob))
```



The dropout rate (e.g. 50%) will define how many nodes are randomly ignored


- Helps to prevent over-fitting
- Speeds up training

Essentially it forces the network to not really too heavily on a single set of neurons.

Code time!



```
# Flatten the data the dropout rate  
model.add(Flatten())
```



..because we are going to feed in a series of fully connected layers which expect a 1D vector

Code time!



```
model.add(Dense(100, activation='elu'))  
model.add(Dense(50, activation='elu'))  
model.add(Dense(10, activation='elu'))  
model.add(Dense(1))
```

Series of fully connected layers
condensing the output into a single value,
i.e. the steering angle.



Dense = fully connected layer

Training the model

- Minimize root mean squared error (RMSE)
 - Between model output and human output
- Adam optimizer (gradient descent)
 - Adaptive moment estimation.

Self Driving



Self Driving

Start the 3D simulation

(double click using mouse)

Ready for instructions

(run in autonomous mode)

Self Driving

Move into source code repository

```
cd {path_to_your_local_repo}
```

Start the car brain

```
python drive.py {model_name.h5}
```

Well Done!



NASCAR Bonus Challenge

Who can drive the fastest?

max speed on line 41 of drive.py

